Differentially Private Set Union

Sivakanth Gopi, Pankaj Gulhane, Janardhan Kulkarni, Judy Hanwen Shen, Milad Shokouhi
and Sergey Yekhanin*

Microsoft
{sigopi,pagulhan,jakul,hashe,milads,yekhanin}@microsoft.com

February 25, 2020

Abstract

We study the basic operation of set union in the global model of differential privacy. In this problem, we are given a universe U of items, possibly of infinite size, and a database D of users. Each user i contributes a subset $W_i \subseteq U$ of items. We want an (ε, δ)-differentially private algorithm which outputs a subset $S \subset \cup_i W_i$ such that the size of S is as large as possible. The problem arises in countless real world applications; it is particularly ubiquitous in natural language processing (NLP) applications as vocabulary extraction. For example, discovering words, sentences, n-grams etc., from private text data belonging to users is an instance of the set union problem.

Known algorithms for this problem proceed by collecting a subset of items from each user, taking the union of such subsets, and disclosing the items whose noisy counts fall above a certain threshold. Crucially, in the above process, the contribution of each individual user is always independent of the items held by other users, resulting in a wasteful aggregation process, where some item counts happen to be way above the threshold. We deviate from the above paradigm by allowing users to contribute their items in a *dependent fashion*, guided by a *policy*. In this new setting ensuring privacy is significantly delicate. We prove that any policy which has certain *contractive* properties would result in a differentially private algorithm. We design two new algorithms, one using Laplace noise and other Gaussian noise, as specific instances of policies satisfying the contractive properties. Our experiments show that the new algorithms significantly outperform previously known mechanisms for the problem.

1 Introduction

Natural language models for applications such as suggested replies for e-mails and dialog systems rely on the discovery of n-grams and sentences [HLLC14, KKR+16, CLB+19, DBS19]. Words and phrases used for training come from individuals, who may be left vulnerable if personal information is revealed. For example, a model could generate a sentence or predict a word that can potentially reveal personal information of the users in the training set [CLE+19]. Therefore, algorithms that allow the public release of the words, n-grams, and sentences obtained from users' text while preserving privacy are desirable. Additional applications of this problem include the release of search queries and keys in SQL queries [KKMN09, WZL+20]. While other privacy definitions are common in practice, guaranteeing differential privacy, introduced in the seminal work of Dwork *et al* [DMNS06], ensures users the strongest preservation of privacy. In this paper we consider user level privacy.

Definition 1.1 (Differential Privacy [DR14a]). A randomized algorithm $\mathcal{A}$ is (ε, δ)-differentially private if for any two neighboring databases D and D', which differ in exactly the data pertaining to a single user, and for all sets $\mathcal{S}$ of possible outputs:

$$\Pr[\mathcal{A}(D) \in \mathcal{S}] \leq e^{\varepsilon} \Pr[\mathcal{A}(D') \in \mathcal{S}] + \delta.$$

---

*Author ordering is alphabetical.

An algorithm satisfying differential privacy (DP) guarantees that its output does not change by much if a single user is either added or removed from the dataset. Moreover, the guarantee holds regardless of how the output of the algorithm is used downstream. Therefore, items (e.g. n-grams) produced using a DP algorithm can be used in other applications without any privacy concerns. Since its introduction a decade ago [DMNS06], differential privacy has become the de facto notion of privacy in statistical analysis and machine learning, with a vast body of research work (see Dwork and Roth [DR+14b] and Vadhan [Vad17] for surveys) and growing acceptance in industry. Differential privacy is deployed in many industries, including Apple [App17], Google [EPK14, BEM+17], Microsoft [DKY17], Mozilla [AKZ+17], and the US Census Bureau [Abo16, KCK+18].

The vocabulary extraction and n-gram discovery problems mentioned above, as well as many commonly studied problems [KKMN09, WZL+20], can be abstracted as a set union which leads to the following problem.

Problem 1.1 (Differentially Private Set Union (DPSU)). Let U be some universe of items, possibly of unbounded size. Suppose we are given a database D of users where each user i has a subset $W_i \subseteq U$. We want an (ε,δ)-differentially private Algorithm A which outputs a subset $S \subseteq \cup_i W_i$ such that the size of S is as large as possible.

Since the universe of items can be unbounded, as in our motivating examples, it is not clear how to apply the exponential mechanism [MT07] to DPSU. Furthermore, even for the cases when U is bounded, implementing the exponential mechanism can be also very inefficient. Existing algorithms [1] for this problem [KKMN09, WZL+20] collect a bounded number of items from each user, build a histogram of these items, and disclose the items whose noisy counts fall above a certain threshold. In these algorithms, the contribution of each user is always *independent* from the identity of items held by other users, resulting in a wasteful aggregation process, where some items' counts could be far above the threshold. Since the goal is to release as large a set as possible rather than to release accurate counts of each item, there could be more efficient ways to allocate the weight to users' items.
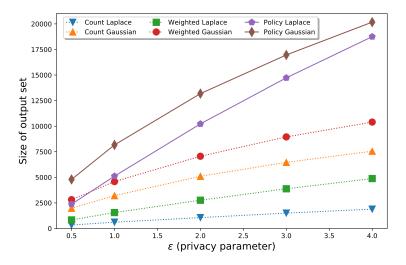

Figure 1: Size of the set output by our proposed algorithms POLICY LAPLACE and POLICY GAUSSIAN compared to natural generalizations of previously known algorithms for various values of privacy parameter ε and $\delta = \exp(-10)$.

We deviate from the previous methods by allowing users to contribute their items in a *dependent fashion*, guided by an *update policy*. In our algorithms, proving privacy is more delicate as some update policies can result in histograms with unbounded sensitivity. We prove a meta-theorem to show that update policies with certain *contractive properties* would result in differentially private algorithms. The main contributions of the paper are:

---

[1]They don't study the DPSU problem as defined in this paper. Their goal is to output approximate counts of as many items as possible in $\cup_i W_i$.

- Guided by our meta-theorems, we introduce two new algorithms called POLICY LAPLACE and POLICY GAUSSIAN for the DPSU problem. Both of them run in *linear time* and only require a single pass over the users' data.

- Using a Reddit dataset, we demonstrate that our algorithms significantly improve the size of DP set union even when compared to natural generalizations of the existing mechanisms for this problem (see Figure 1).

Our algorithms are being productized in industry to make a basic subroutine in an NLP application differentially private.

1.1 Baseline algorithms

To understand the DPSU problem better, let us start with the simplest case we can solve by known techniques. Define $\Delta_0 = \max_i |W_i|$. Suppose $\Delta_0 = 1$. This special case can be solved using the algorithms in [KKMN09, WZL$^+$20]. Their algorithm works as follows: Construct a histogram on $\cup_i W_i$ (the set of items in a database D) where the count of each item is the number of sets it belongs to. Then add Laplace noise or Gaussian noise to the counts of each item. Finally, release only those items whose noisy histogram counts are above a certain *threshold* ρ. It is not hard to prove that if the threshold is set sufficiently high, then the algorithm is (ε, δ)-DP.

A straight-forward extension of the histogram algorithm for $\Delta_0 > 1$ is to upper bound the ℓ_1-sensitivity by Δ_0 (and ℓ_2-sensitivity by $\sqrt{\Delta_0}$), and then add some appropriate amount of Laplace noise (or Gaussian noise) based on sensitivity. The threshold ρ has to be set based on Δ_0. The Laplace noise based algorithm was also the approach considered in [KKMN09, WZL$^+$20]. This approach has the following drawback. Suppose a significant fraction of users have sets of size smaller than Δ_0. Then constructing a histogram based on *counts* of the items results in *wastage of sensitivity budget*. A user i with $|W_i| < \Delta_0$ can increment the count of items in W_i by any vector $v \in \mathbb{R}^{W_i}$ as long as one can ensure that ℓ_1 sensitivity is bounded by Δ_0 (or ℓ_2 sensitivity is bounded by $\sqrt{\Delta_0}$ if adding Gaussian noise). Consider the following natural generalization of Laplace and Gaussian mechanisms to create a *weighted histogram of elements*. A weighted histogram over a domain X is any map $H : X \to \mathbb{R}$. For an item $u \in U$, $H(u)$ is called the weight of u. In the rest of the paper, the term histogram should be interpreted as weighted histogram. Each user i updates the weight of each item $u \in W_i$ using the rule: $H[u] := H[u] + (\Delta_0/|W_i|)^{1/p}$ for $p = 1, 2$. It is not hard to see that ℓ_p-sensitivity of this weighted histogram is still $\Delta_0^{1/p}$. Adding Laplace noise (for $p = 1$) or Gaussian noise (for $p = 2$) to each item of the weighted histogram, and releasing only those items above an appropriately calibrated threshold will lead to differentially private output. We call these algorithms as WEIGHTED LAPLACE and WEIGHTED GAUSSIAN, they will be used as benchmarks to compare against our new algorithms.

1.2 Our techniques

The WEIGHTED LAPLACE and WEIGHTED GAUSSIAN mechanisms described above can be thought of trying to solve the following variant of a *Knapsack* problem. Here each item $u \in U$ is a bin and we gain a profit of 1 if the total weight of the item in the weighted histogram constructed is more than the threshold. Each user can increment the weight of elements $u \in W_i$ using an *update policy* ϕ which is defined as follows.

Definition 1.2 (Update policy). An update policy is a map $\phi : \mathbb{R}^U \times 2^U \to \mathbb{R}^U$ such that $\text{supp}(\phi(H, W) - H) \subset W$, i.e., ϕ can only update the weights of items in W. And the i^{th} user updates H to $\phi(H, W_i)$. Since W_i is typically understood from context, we will write $\phi(H)$ instead of $\phi(H, W_i)$ for simplicity.

In this framework, the main technical challenge is the following:

How to design update policies such that the sensitivity of the resulting weighted histogram is small while maximizing the number of bins that are full?

Note that bounding sensitivity requires that $\|\phi(H, W) - H\|_{\ell_p} \leq C$ for some constant C i.e. each user has an ℓ_p-budget of C and can increase the weights of items in their set by an ℓ_p-distance of at most C. By scaling, WLOG we can assume that $C = 1$. Note that having a larger value of Δ_0 should help in filling more bins as users have more choice in how they can use their budget to increment the weight of items.

In this paper, we consider algorithms which *iteratively* construct the weighted histogram. That is, in our algorithms, we consider users in a random order, and each user updates the weighted histogram using the update policy ϕ. Algorithm 1 is a meta-algorithm for DP set union, and all our subsequent algorithms follow this framework.

Algorithm 1 High level meta algorithm for DP Set Union

Input: D: Database of n users where each user i has some subset $W_i \subset U$
ρ: threshold
Noise: Noise distribution ($\mathsf{Lap}(0, \lambda)$ or $\mathcal{N}(0, \sigma^2)$)
Output: S: A subset of $\cup_i W_i$
Build weighted histogram H supported over $\cup_i W_i$ using Algorithm 2.
$S = \{\}$ (empty set)
for $u \in \cup_i W_i$ **do**
   $\hat{H}[u] \leftarrow H[u] + \mathsf{Noise}$
   if $\hat{H}[u] > \rho$ **then**
      $S \leftarrow S \cup \{u\}$
   end if
end for
Output S

Algorithm 2 High level meta algorithm for building weighted histogram using a given update policy

---

Input: D: Database of n users where each user i has some subset $W_i \subset U$
Δ_0: maximum contribution parameter
hash: A random hash function which maps user ids into some large domain without collisions
ϕ: Update policy for a user to update the weights of items in their set
Output: H: A weighted histogram in $\mathbb{R}^{\cup_i W_i}$
$H = \{\}$ (empty histogram)
Sort users into $\mathrm{User}_1, \mathrm{User}_2, \ldots, \mathrm{User}_n$ by sorting the hash values of their user ids under the hash function hash
for $i = 1$ **to** n **do**
   $W_i \leftarrow$ set with User_i
   if $|W_i| > \Delta_0$ **then**
      $W_i' \leftarrow$ Randomly choose Δ_0 items from W_i
   else
      $W_i' \leftarrow W_i$
   end if
   Update $H[u]$ for each $u \in W_i'$ using update policy ϕ
end for
Output H

If the update policy is such that it increments the weights of items independent of other users (as done in WEIGHTED LAPLACE and WEIGHTED GAUSSIAN), then it is not hard to see that sensitivity of H can be bounded by 1; that is, by the budget of each user. However, if some item is already way above the threshold ρ, then it does not make much sense to waste the limited budget on that item. Ideally, users can choose a clever *update policy* to distribute their budget among the W_i items based on the current weights.

Note that if a policy is such that updates of a user depends on other users, it can be quite tricky to bound the sensitivity of the resulting weighted histogram. To illustrate this, consider for example the *greedy* update policy. Each user i can use his budget of 1 to fill the bins that is closest to the threshold among the bins $u \in W_i$. If an item already reached the threshold, the user can spend his remaining budget incrementing the weight of next bin that is closest to the threshold and so on. Note that from our Knapsack problem analogy this seems be a good way to maximize the number of bins filled. However such a greedy policy can have very large sensitivity (see appendix for an example), and hence won't lead to any reasonable DP algorithm. So,

the main contribution of the paper is in showing policies which help maximize the number of item bins that are filled while keeping the sensitivity low. In particular, we define a general class of ℓ_p-contractive update policies and show that they produce weighted histograms with bounded ℓ_p-sensitivity.

Definition 1.3 (ℓ_p-contractive update policy)**.** We say that an update policy ϕ is ℓ_p-contractive if there exists a subset $\mathcal{I}$ (called the invariant subset for ϕ) of pairs of weighted histograms which are at an ℓ_p distance of at most 1, i.e.,

$$\mathcal{I} \subset \left\{ (H_1, H_2) : \|H_1 - H_2\|_{\ell_p} \leq 1 \right\}$$

such that the following conditions hold.

1. (Invariance) $(H_1, H_2) \in \mathcal{I} \Rightarrow (\phi(H_1), \phi(H_2)) \in \mathcal{I}$.[2]

2. $(\phi(H), H) \in \mathcal{I}$ for all H.

Property (2) of Definition 1.3 requires that the update policy can change the histogram by an ℓ_p distance of at most 1 (budget of a user).

Theorem 1.1 (Contractivity implies bounded sensitivity)**.** Suppose ϕ is an update policy which is ℓ_p-contractive over some invariant subset $\mathcal{I}$. Then the histogram output by Algorithm 2 has ℓ_p-sensitivity bounded by 1.

We prove Theorem 1.1 in Section 3. Once we have bounded ℓ_p-sensitivity, we can get a DP Set Union algorithm with some additional technical work.

Theorem 1.2. (Informal: Bounded sensitivity implies DP) For $p \in \{1, 2\}$, if the ℓ_p-sensitivity of the weighted histogram output by Algorithm 2 is bounded, then Algorithm 1 for DP Set Union can be made (ε, δ)-differentially private by appropriately choosing the noise distribution (Noise) and threshold (ρ).

The main contribution of the paper is two new algorithms guided by Theorem 1.1. The first algorithm, which we call POLICY LAPLACE, uses update policy that is ℓ_1-contractive. The second algorithm, which we call POLICY GAUSSIAN, uses update policy that is ℓ_2-contractive. Finally we show that our algorithms significantly outperform the weighted update policies.

At a very high-level, the role of contractivity in our algorithms is indeed similar to its role in the recent elegant work of Feldman *et al* [FMTT18]. They show that if an iterative algorithm is contractive in each step, then adding Gaussian noise in each iteration will lead to strong privacy amplification. In particular, users who make updates early on will enjoy much better privacy guarantees. However their framework is not applicable in our setting, because their algorithm requires adding noise to the count of every item in every iteration; this will lead to unbounded growth of counts and items which belong to only a single user can also get output which violates privacy.

2 Preliminaries

Let $\mathcal{D}$ denote the collection of all databases. We say that D, D' are neighboring databases, denoted by $D \sim D'$, if they differ in exactly one user.

Definition 2.1. For $p \geq 0$, the ℓ_p-sensitivity of $f : \mathcal{D} \to \mathbb{R}^k$ is defined as $\sup_{D \sim D'} \|f(D) - f(D')\|_{\ell_p}$ where the supremum is over all neighboring databases D, D'.

Proposition 2.1 (The Laplace Mechanism [DR14a])**.** Given any function $f : D \to \mathbb{R}^k$, the Laplace Mechanism is defined as:

$$\mathcal{M}(x, f(.), \varepsilon) = f(x) + (Y_i, ..., Y_k) \tag{1}$$

where Δ_1 is the ℓ_1-sensitivity and Y_i are i.i.d. random variables drawn from $\mathsf{Lap}(0, \Delta_1/\varepsilon)$.

---

[2]Note that property (1) is a slightly weaker requirement than the usual notion of ℓ_p-contractivity which requires $\|\phi(H_1) - \phi(H_2)\|_{\ell_p} \leq \|H_1 - H_2\|_{\ell_p}$ for all H_1, H_2. Instead we require contraction only for $(H_1, H_2) \in \mathcal{I}$.

Proposition 2.2 (Gaussian Mechanism [BW18]). *If $f : \mathcal{D} \to \mathbb{R}^d$ is a function with ℓ_2-sensitivity Δ_2. For any $\varepsilon \geq 0$ and $\delta \in [0, 1]$, the Gaussian output perturbation mechanism $M(x) = f(x) + Z$ with $Z \sim \mathcal{N}(0, \sigma^2 I)$ is (ε, δ)-DP if and only if*

$$\Phi\left(\frac{\Delta_2}{2\sigma} - \frac{\varepsilon\sigma}{\Delta_2}\right) - e^{\varepsilon}\Phi\left(-\frac{\Delta_2}{2\sigma} - \frac{\varepsilon\sigma}{\Delta_2}\right) \leq \delta.$$

Definition 2.2. We say that two distributions P, Q on a domain X are (ε, δ)-close to each other, denoted by $P \approx_{\varepsilon,\delta} Q$, if for every $S \subset X$, we have

1. $\Pr_{x \sim P}[x \in S] \leq e^{\varepsilon} \Pr_{x \sim Q}[x \in S] + \delta$ and

2. $\Pr_{x \sim Q}[x \in S] \leq e^{\varepsilon} \Pr_{x \sim P}[x \in S] + \delta$.

We say that two random variables X, Y are (ε, δ)-close to each other, denoted by $X \approx_{\varepsilon,\delta} Y$, if their distributions are (ε, δ)-close to each other.

We will need the following lemmas which are useful to prove (ε, δ)-DP.

Lemma 2.1. *Let P, Q be probability distributions over a domain X. If there exists an event E s.t. $P(E) = 1 - \delta'$ and $P|_E \approx_{\varepsilon,\delta} Q$, then $P \approx_{\varepsilon,\delta+\delta'} Q$.*

Proof. Fix some subset $S \subseteq X$.

$$\Pr_{x \sim P}[x \in S] = P(\bar{E}) \Pr_{x \sim P}[x \in S|\bar{E}] + P[E] \Pr_{x \sim P}[x \in S|E]$$
$$\leq P(\bar{E}) + \Pr_{x \sim P}[x \in S|E]$$
$$= \delta' + \Pr_{x \sim P|_E}[x \in S]$$
$$\leq \delta' + e^{\varepsilon} \Pr_{x \sim Q}[x \in S] + \delta$$

We now prove the other direction.

$$\Pr_{x \sim Q}[x \in S] \leq e^{\varepsilon} \Pr_{x \sim P|_E}[x \in S] + \delta$$
$$\leq e^{\varepsilon} \frac{\Pr_{x \sim P}[x \in S]}{P(E)} + \delta$$
$$= e^{\varepsilon} \frac{\Pr_{x \sim P}[x \in S]}{1 - \delta'} + \delta$$
$$= e^{\varepsilon} \Pr_{x \sim P}[x \in S] + \delta'\left(\frac{e^{\varepsilon} \Pr_{x \sim P}[x \in S]}{1 - \delta'}\right) + \delta$$

Now if $e^{\varepsilon} \Pr_{x \sim P}[x \in S] \leq 1 - \delta'$, then we have $Pr_{x \sim Q}[x \in S] \leq e^{\varepsilon} \Pr_{x \sim P}[x \in S] + \delta' + \delta$. Otherwise, trivially

$$\Pr_{x \sim Q}[x \in S] \leq 1 \leq e^{\varepsilon} \Pr_{x \sim P}[x \in S] + \delta' + \delta.$$

$\square$

We will also need the fact that if $X \approx_{\varepsilon,\delta} Y$, then after post-processing they also remain (ε, δ)-close.

Lemma 2.2. *If two random variables X, Y are (ε, δ)-close and M is any randomized algorithm, then $M(X) \approx_{\varepsilon,\delta} M(Y)$.*

Proof. Let $M(z) = F(z, R)$ for some function F where R is the random bits used by M. For any subset S of the possible outputs of M,

$$
\begin{aligned}
\Pr[M(X) \in S] &= \Pr_{X,R}[F(X, R) \in S] \\
&= \sum_r \Pr[R = r] \Pr_X[F(X, r) \in S] \\
&\leq \sum_r \Pr[R = r] \left(e^\varepsilon \Pr_Y[F(Y, r) \in S] + \delta \right) \\
&= e^\varepsilon \sum_r \Pr[R = r] \Pr_Y[F(Y, r) \in S] + \delta \\
&= e^\varepsilon \Pr_{X,R}[F(Y, R) \in S] + \delta.
\end{aligned}
$$

The other direction holds by symmetry. $\square$

3 Contractivity implies DP algorithms

In this section, we show that if an update policy satisfies contractive property as in Definition 1.3, we can use it to develop a DPSU algorithm. First we show that contractivity of update policy implies bounded sensitivity (Theorem 1.1), which in turn implies a DPSU algorithm by Theorem 1.2. We will first define sensitivity and update policy formally. Let $\mathcal{D}$ denote the collection of all databases. We say that D, D' are neighboring databases, denoted by $D \sim D'$, if they differ in exactly one user.

Definition 3.1. For $p \geq 0$, the ℓ_p-sensitivity of $f : \mathcal{D} \to \mathbb{R}^k$ is defined as $\sup_{D \sim D'} \|f(D) - f(D')\|_{\ell_p}$ where the supremum is over all neighboring databases D, D'.

Proof of Theorem 1.1. Let ϕ be an ℓ_p-contractive update policy with invariant subset $\mathcal{I}$. Consider two neighboring databases D_1 and D_2 where D_1 has one extra user compared to D_2. Let H_1 and H_2 denote the histograms built by Algorithm 1 using the update policy ϕ when the databases are D_1 and D_2 respectively.

Say the extra user in D_1 has position t in the global ordering given by the hash function. Let H_1^{t-1} and H_2^{t-1} be the histograms after the first $t-1$ (according to the global order given by the hash function hash) users' data is added to the histogram. Therefore $H_1^{t-1} = H_2^{t-1}$. And the new user updates H_1^{t-1} to H_1^t. By property (2) in Definition 1.3 of ℓ_p-contractive policy, $(\phi(H_1^{t-1})), H_1^{t-1}) \in \mathcal{I}$. Since $\phi(H_1^{t-1}) = H_1^t$, we have $(H_1^t, H_1^{t-1}) = (H_1^t, H_2^{t-1}) \in \mathcal{I}$. The remaining users are now added to H_1^t, H_2^{t-1} in the same order. Note that we are using the fact that the users are sorted according some hash function and they contribute in that order (this is also needed to claim that $H_1^{t-1} = H_2^{t-1}$). Therefore, by property (1) in Definition 1.3 of ℓ_p-contractive policy, we get $(H_1, H_2) \in \mathcal{I}$. Since $\mathcal{I}$ only contains pairs with ℓ_p-distance at most 1, we have $\|H_1 - H_2\|_{\ell_p} \leq 1$. Therefore the histogram built by Algorithm 2 using ϕ has ℓ_p-sensitivity of at most 1. $\square$

Above theorem implies that once we have a ℓ_p contractive update policy, we can appeal to Theorem 1.2 to design an algorithm for DPSU.

4 Policy Laplace algorithm

In this section we will present an ℓ_1-contractive update policy called ℓ_1-DESCENT (Algorithm 3) and use it to obtain a DP Set Union algorithm called POLICY LAPLACE (Algorithm 4).

4.1 ℓ_1-descent update policy

The policy is described in Algorithm 3. We will set some *cutoff* Γ above the threshold ρ to use in the update policy. Once the weight of an item ($H[u]$) crosses the cutoff, we do not want to increase it further. In this policy, each user starts with a budget of 1. The user uniformly increases $H[u]$ for each $u \in W_i'$ s.t. $H[u] < \Gamma$. Once some item's weight reaches Γ, the user stops increasing that item and keeps increasing the rest of the

items until the budget of 1 is expended. To implement this efficiently, the Δ_0 items from each user are sorted based on distance to the cutoff. Beginning with the item whose weight is closest to the cutoff Γ (but still below the cutoff), say item u, we will add $\Gamma - H[u]$ (gap to cutoff for item u) to each of the items below the cutoff. This repeats until the user's budget of 1 has been expended.

This policy can also be interpreted as *gradient descent* to minimize the ℓ_1-distance between the current weighted histogram and the point $(\Gamma, \Gamma, \dots, \Gamma)$, hence the name ℓ_1-DESCENT. Since the gradient vector is 1 in coordinates where the weight is below cutoff Γ and 0 in coordinates where the weight is Γ, the ℓ_1-DESCENT policy is moving in the direction of the gradient until it has moved a total ℓ_1-distance of at most 1.

Algorithm 3 ℓ_1-DESCENT update policy

  Input: H: Current histogram
  W: A subset of U of size at most Δ_0
  Γ: cutoff parameter
  Output: H: Updated histogram
  // Build cost dictionary G
  $G = \{\}$ // Empty dictionary
  for $u \in W$ **do**
    if $H[u] < \Gamma$ **then**
      // Gap to cutoff for items below cutoff Γ
      $G[u] \leftarrow \Gamma - H[u]$
    end if
  end for
  budget $\leftarrow 1$ // Each user gets a total budget of 1
  $K \leftarrow |G|$ // Number of items still under cutoff
  // Sort in increasing order of the gap $\Gamma - H[u]$
  $G \leftarrow \text{sort}(G)$
  // Let $u_1, u_2, \dots, u_{|B|}$ be the sorted order
  for $j = 1$ to $|B|$ **do**
    // Cost of increasing weights of remaining K items by $G[u_j]$
    cost $= G[u_j] \cdot K$
    if cost $\leq$ budget **then**
      for $\ell = j$ to $|G|$ **do**
        $H[u_\ell] \leftarrow H[u_\ell] + G[u_j]$
        // Gap to cutoff is reduced by $G[u_j]$
        $G[u_\ell] \leftarrow G[u_\ell] - G[u_j]$
      end for
      budget $\leftarrow$ budget - cost
      // u_j has reached cutoff, so decrease K by 1
      $K \leftarrow K - 1$
    else
      for $\ell = j$ to $|G|$ **do**
        // Update item weights by as much as remaining budget allows
        $H[u_\ell] \leftarrow H[u_\ell] + \frac{\text{budget}}{K}$
      break
      end for
    end if
  end for

4.2 Policy Laplace

The POLICY LAPLACE algorithm (Algorithm 4) for DPSU uses the framework of the meta algorithm in Algorithm 1 with the update policy in Algorithm 3. Since the added noise is $\mathsf{Lap}(0, \lambda)$, which is centered at 0, we want to set the cutoff Γ in the update policy to be sufficiently above the threshold ρ. Thus we pick

$\Gamma = \rho_{\mathsf{Lap}} + \alpha \cdot \lambda$ for some $\alpha > 0$. From our experiments, choosing $\alpha \in [2, 6]$ works best empirically. The parameters $\lambda, \rho_{\mathsf{Lap}}$ are set so as to achieve (ε, δ)-DP as shown in Theorem 4.1.

Algorithm 4 POLICY LAPLACE algorithm for DPSU

Input: D: Database of n users where each user has some subset $W \subset U$
Δ_0: maximum contribution parameter
(ε, δ): privacy parameters
α: parameter for setting cutoff
Output: S: A subset of $\cup_i W_i$
$\lambda \leftarrow 1/\varepsilon$ // Noise parameter in $\mathsf{Lap}(0, \lambda)$
// Threshold parameter
$\rho_{\mathsf{Lap}} \leftarrow \max_{1 \leq t \leq \Delta_0} \frac{1}{t} + \frac{1}{\varepsilon} \log\left(\frac{1}{2\left(1-(1-\delta)^{1/t}\right)}\right)$
$\Gamma \leftarrow \rho_{\mathsf{Lap}} + \alpha \cdot \lambda$ // Cutoff parameter for update policy
Run Algorithm 1 with $\mathsf{Noise} \sim \mathsf{Lap}(0, \lambda)$ and the ℓ_1-DESCENT update policy in Algorithm 3 to output S.

4.3 Privacy analysis of Policy Laplace

In this section, we will prove that the POLICY LAPLACE algorithm (Algorithm 4) is (ε, δ)-DP. By Theorem 1.2 and Theorem 1.1, it is enough to show that ℓ_1-DESCENT policy (Algorithm 3) is ℓ_1-contractive. For two histograms G_1, G_2, we write $G_1 \geq G_2$ if $G_1[u] \geq G_2[u]$ for each every item u. $G_1 \leq G_2$ is defined similarly.

Lemma 4.1. Let $\mathcal{I} = \{(G_1, G_2) : G_1 \geq G_2, \|G_1 - G_2\|_{\ell_1} \leq 1\}$. Then ℓ_1-DESCENT update policy is ℓ_1-contractive over the invariant subset $\mathcal{I}$.

Proof. Let ϕ denote the ℓ_1-DESCENT update policy.

We will first show property (2) of Definition 1.3. Let G be any weighted histogram and let $G' = \phi(G)$. Clearly $G' \geq G$ as the new user will never decrease the weight of any item. Moreover, the total change to the histogram is at most 1 in ℓ_1-distance. Therefore $\|G' - G\|_{\ell_1} \leq 1$. Therefore $(G', G) \in \mathcal{I}$.

We will now prove property (1) of Definition 1.3. Let $(G_1, G_2) \in \mathcal{I}$, i.e., $G_1 \geq G_2$ and $\|G_1 - G_2\|_{\ell_1} \leq 1$. Let $G_1' = \phi(G_1), G_2' = \phi(G_2)$. A new user can increase G_1 and G_2 by at most 1 in ℓ_1 distance. Let Γ be the cutoff parameter in Algorithm 3. Let S be the set of Δ_0 items with the new user, therefore only the items in S will change in G_1', G_2'. WLOG, we can assume that the user changes both G_1 and G_2 by exactly total ℓ_1 distance of 1. Otherwise, in at least one of them all the items in S should reach the cutoff Γ. If this happens with G_1, then clearly $\Gamma = G_1'[u] \geq G_2'[u]$ for all $u \in S$. But it is easy to see that if this happens with G_2, then it should also happen with G_1 in which case $G_1'[u] = G_2'[u] = \Gamma$ for $u \in S$.

Imagine that at time $t = 0$, the user starts pushing mass continuously at a rate of 1 to both G_1, G_2 until the entire mass of 1 is sent, which happens at time $t = 1$. The mass flow is equally split among all the items which haven't yet crossed cutoff. Let G_1^t and G_2^t be the histograms at time t. Therefore $G_i^0 = G_i$ and $G_i^1 = G_i'$. We claim that $G_1^t \geq G_2^t$ implies that $\frac{dG_1^t[u]}{dt} \geq \frac{dG_2^t[u]}{dt}$ for all $u \in S$ s.t. $G_1^t[u] < \Gamma$. This is because the flow is split equally among items which didn't cross the cutoff, and there can only be more items in G_2^t which didn't cross the the cutoff when compared to G_1^t. And at time $t = 0$, we have $G_1^0 \geq G_2^0$. Therefore, we have $G_1^t \geq G_2^t$ for all $t \in [0, 1]$ and so $G_1' \geq G_2'$.

We will now prove ℓ_1-contraction. Let $C_i = \|G_i - G_i'\|_{\ell_1}$. By the discussion above, $C_1 \leq C_2$ (either total

mass flow is equal to 1 for both or all items in S will reach cutoff Γ in G_1 before this happens in G_2).

$$
\begin{aligned}
\|G_1' &- G_2'\|_{\ell_1} \\
&= \sum_{u \in S} G_1'[u] - \sum_{u \in S} G_2'[u] && \text{(Since } G_1' \geq G_2') \\
&= \sum_{u \in S} G_1[u] - \sum_{u \in S} G_2[u] + C_1 - C_2 \\
&\leq \sum_{u \in S} G_1[u] - \sum_{u \in S} G_2[u] && \text{(Since } C_1 \leq C_2) \\
&= \|G_1 - G_2\|_{\ell_1} && \text{(Since } G_1 \geq G_2) \\
&\leq 1.
\end{aligned}
$$

Therefore $(G_1', G_2') \in \mathcal{I}$ which proves property (2) of Definition 1.3. □

Lemma 4.2. Suppose D_1 and D_2 are neighboring databases where D_1 has one extra user compared to D_2. Let H_1 and H_2 denote the histograms built by the POLICY LAPLACE algorithm (Algorithm 4) when the database is D_1 and D_2 respectively. Then $\|H_1 - H_2\|_{\ell_1} \leq 1$.

Proof. Say the extra user in D_1 has position t in the global ordering given by the hash function. Let H_1^i and H_2^i be the histograms after the first i (according to the global order) users data is added to the histogram. Therefore $H_1^{t-1} = H_2^{t-1}$. And the new user updates H_1^{t-1} to H_1^t. Since the total ℓ_1 change by an user is at most 1, $\|H_1^t - H_1^{t-1}\|_{\ell_1} = \|H_1^t - H_2^{t-1}\|_{\ell_1} \leq 1$. The remaining users are now added to H_1^t, H_2^{t-1} in the same order. Note that we are using the fact that the users are sorted according some hash function and they contribute in that order (this is also needed to claim that $H_1^{t-1} = H_2^{t-1}$). Therefore, by the ℓ_1-contraction property shown in Lemma 4.1, $\|H_1 - H_2\|_{\ell_1} \leq 1$. □

We now state a formal theorem which proves $(\varepsilon, \delta) - DP$ of POLICY LAPLACE algorithm.

Theorem 4.1. The POLICY LAPLACE algorithm (Algorithm 4) is (ε, δ)-DP when

$$
\rho_{\text{Lap}} \geq \max_{1 \leq t \leq \Delta_0} \frac{1}{t} + \frac{1}{\varepsilon} \log \left(\frac{1}{2 \left(1 - (1 - \delta)^{1/t} \right)} \right).
$$

Proof. Suppose D_1 and D_2 are neighboring databases where D_1 has one extra user compared to D_2. Let P and Q denote the distribution of output of the algorithm when the database is D_1 and D_2 respectively. We want to show that $P \approx_{\varepsilon, \delta} Q$. Let E be the event that $A \subset \text{supp}(H_2)$.

Claim 4.1. $P|_E \approx_{\varepsilon, 0} Q$

Proof. Let H_1 and H_2 be the histograms generated by the algorithm from databases D_1 and D_2 respectively. And $\hat{H}_1$ and $\hat{H}_2$ be the histograms obtained by adding $\text{Lap}(0, 1/\varepsilon)$ noise to each entry of H_1 and H_2 respectively. For any possible output A of Algorithm 4, we have

$$
Q(A) = \Pr[A = \{u \in \text{supp}(H_2) : \hat{H}_2[u] > \rho_{\text{Lap}}\}] \text{ and } P|_E(A) = \Pr[A = \{u \in \text{supp}(H_2) : \hat{H}_1[u] > \rho_{\text{Lap}}\}].
$$

So $A \sim P|_E$ is obtained by post-processing $\hat{H}_1|_E$ and $A \sim Q$ is obtained by post-processing $\hat{H}_2$. Since post-processing only makes two distributions closer (Lemma 2.2), it is enough to show that the distributions of the $\hat{H}_1|_{\text{supp}(H_2)}$ and $\hat{H}_2$ are $(\varepsilon, 0)$-close to each other. By Lemma 4.2, $H_1|_{\text{supp}(H_2)}$ and H_2 differ in ℓ_1-distance by at most 1. Therefore $P|_E \approx_{\varepsilon, 0} Q$ by the properties of Laplace mechanism (see Theorem 3.6 in [DR14a]). □

By Lemma 2.1, it is enough to show that $P(E) \geq 1 - \delta$. Let $T = \text{supp}(H_1) \setminus \text{supp}(H_2)$. Note that $|T| \leq \Delta_0$ and $H_1[u] \leq \frac{1}{|T|}$ for $u \in T$.

$$P(\bar{E}) = \Pr[\exists u \in T \mid \hat{H}_1[u] > \rho_{\mathsf{Lap}}]$$

$$= 1 - \Pr[\forall u \in T \quad \hat{H}_1[u] \leq \rho_{\mathsf{Lap}}]$$

$$= 1 - \prod_{u \in T} \Pr[H_1[u] + X_u \leq \rho_{\mathsf{Lap}}] \qquad (X_u \sim \mathsf{Lap}(1/\varepsilon))$$

$$\leq 1 - \prod_{u \in T} \Pr\left[X_u \leq \rho_{\mathsf{Lap}} - \frac{1}{|T|}\right] \qquad (H_1[u] \leq \tfrac{1}{|T|} \text{ for } u \in T)$$

$$= 1 - \left(1 - \frac{1}{2} \exp\left(-\varepsilon \rho_{\mathsf{Lap}} + \varepsilon \frac{1}{|T|}\right)\right)^{|T|} \qquad (2)$$

Thus for

$$\rho_{\mathsf{Lap}} \geq \max_{1 \leq t \leq \Delta_0} \frac{1}{t} + \frac{1}{\varepsilon} \log\left(\frac{1}{2\left(1 - (1-\delta)^{1/t}\right)}\right),$$

we have $P(\bar{E}) \leq \delta$. Therefore the POLICY LAPLACE algorithm (Algorithm 4) is (ε, δ)-DP. $\qquad \square$

5 Policy Gaussian algorithm

In this section we will present an ℓ_2-contractive update policy called ℓ_2-DESCENT (Algorithm 5) and use it to obtain a DP Set Union algorithm called POLICY GAUSSIAN (Algorithm 4).

5.1 ℓ_2-descent update policy

Similar to the Laplace update policy, we will set some *cutoff* Γ above the threshold ρ and once an item's count ($H[u]$) crosses the cutoff, we don't want to increase it further. In this policy, each user starts with a budget of 1. But now, the total change a user can make to the histogram can be at most 1 when measured in ℓ_2-norm (whereas in Laplace update policy we used ℓ_1-norm to measure change). In other words, sum of the squares of the changes that the user makes is at most 1. Since we want to get as close to the cutoff (Γ) as possible, the user moves the counts vector (restricted to the set W of Δ_0 items the user has) in the direction of the point $(\Gamma, \Gamma, \ldots, \Gamma)$ by an ℓ_2-distance of at most 1. This update policy is presented in Algorithm 5.

This policy can also be interpreted as *gradient descent* to minimize the ℓ_2-distance between the current weighted histogram and the point $(\Gamma, \Gamma, \ldots, \Gamma)$, hence the name ℓ_2-DESCENT. Since the gradient vector is in the direction of the line joining the current point and $(\Gamma, \Gamma, \ldots, \Gamma)$, the ℓ_2-DESCENT policy is moving the current histogram towards $(\Gamma, \Gamma, \ldots, \Gamma)$ by an ℓ_2-distance of at most 1.

5.2 Policy Gaussian

The POLICY GAUSSIAN algorithm (Algorithm 6) for DPSU uses the framework of the meta algorithm in Algorithm 1 using the Gaussian update policy (Algorithm 5). Since the added noise is $\mathcal{N}(0, \sigma^2)$ which is centered at 0, we want to set the cutoff Γ in the update policy to be sufficiently above (but not too high above) the threshold ρ_{Gauss}. Thus we pick $\Gamma = \rho_{\mathsf{Gauss}} + \alpha \cdot \sigma$ for some $\alpha > 0$. From our experiments, choosing $\alpha \in [2, 6]$ empirically yields these best results. The parameters $\sigma, \rho_{\mathsf{Gauss}}$ are set so as to achieve (ε, δ)-DP as shown in Theorem 5.1. $\Phi(\cdot)$ is the cumulative density function of standard Gaussian distribution and $\Phi^{-1}(\cdot)$ is its inverse.

To find $\min\left\{\sigma : \Phi\left(\frac{1}{2\sigma} - \varepsilon\sigma\right) - e^{\varepsilon}\Phi\left(-\frac{1}{2\sigma} - \varepsilon\sigma\right) \leq \frac{\delta}{2}\right\}$, one can use binary search because $\Phi\left(\frac{1}{2\sigma} - \varepsilon\sigma\right) - e^{\varepsilon}\Phi\left(-\frac{1}{2\sigma} - \varepsilon\sigma\right)$ is a decreasing function of σ. An efficient and robust implementation of this binary search can be found in [BW18].

Algorithm 5 ℓ_2-DESCENT update policy

Input: H: Current histogram
W: A subset of U of size at most Δ_0
Γ: cutoff parameter
Output: H: Updated histogram
$G = \{\}$ // Empty dictionary
for $u \in W$ **do**
    // G is the vector joining $H|_W$ to $(\Gamma, \Gamma, \ldots, \Gamma)$
    $G[u] \leftarrow \Gamma - H[u]$
end for
// ℓ_2-distance between $H|_W$ and $(\Gamma, \Gamma, \ldots, \Gamma)$
$Z \leftarrow \left(\sum_{u \in W} G[u]^2 \right)^{1/2}$
// If $Z \leq 1$, then the user moves $H|_W$ to $(\Gamma, \Gamma, \ldots, \Gamma)$. Else, move $H|_W$ in the direction of $(\Gamma, \Gamma, \ldots, \Gamma)$ by
an ℓ_2-distance of at most 1
if $Z < 1$ **then**
    for $u \in W$ **do**
        $H[u] \leftarrow \Gamma$
    end for
else
    for $u \in W$ **do**
        $H[u] \leftarrow H[u] + \frac{G[u]}{Z}$
    end for
end if

5.3   Privacy analysis of Policy Gaussian

In this section we will prove that the POLICY GAUSSIAN algorithm (Algorithm 6) is (ε, δ)-DP. By Theorem 1.2 and Theorem 1.1, it is enough to show ℓ_2-contractivity of ℓ_2-DESCENT update policy. We will need a simple plane geometry lemma for this.

Lemma 5.1. Let A, B, C denote the vertices of a triangle in the Euclidean plane. If $|AB| > 1$, let B' be the point on the side AB which is at a distance of 1 from B and if $|AB| \leq 1$, define $B' = A$. C' is defined similarly. Then $|B'C'| \leq |BC|$.
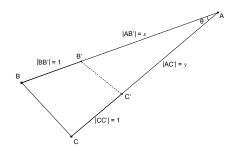

Figure 2: Geometric illustration of Lemma 5.1 when $|AB|, |AC| > 1$. The lemma implies that $|B'C'| \leq |BC|$.

Proof of Lemma 5.1. Let us first assume that both $|AB|, |AC| > 1$. Let θ be the angle at A and let $|AB'| =$

Algorithm 6 POLICY GAUSSIAN algorithm for DPSU
___
Input: D: Database of n users where each user has some subset $W \subset U$
Δ_0: maximum contribution parameter
(ε, δ): privacy parameters
α: parameter for setting cutoff
Output: S: A subset of $\cup_i W_i$
// Standard deviation in Gaussian noise
$\sigma \leftarrow \min \left\{ \sigma : \Phi \left(\frac{1}{2\sigma} - \varepsilon \sigma \right) - e^\varepsilon \Phi \left(-\frac{1}{2\sigma} - \varepsilon \sigma \right) \leq \frac{\delta}{2} \right\}$
// Threshold parameter
$\rho_{\mathsf{Gauss}} \leftarrow \max_{1 \leq t \leq \Delta_0} \left(\frac{1}{\sqrt{t}} + \sigma \Phi^{-1} \left(\left(1 - \frac{\delta}{2} \right)^{1/t} \right) \right)$
$\Gamma \leftarrow \rho_{\mathsf{Lap}} + \alpha \cdot \sigma$ // Cutoff parameter for update policy
Run Algorithm 1 with Noise $\sim \mathcal{N}(0, \sigma^2)$ and the ℓ_2-DESCENT update policy in Algorithm 5 to output S.
___

$x, |AC'| = y$ as shown in Figure 2. Then by the cosine formula,

$$
\begin{aligned}
|BC|^2 &= |AB|^2 + |AC|^2 - 2|AB||AC| \cos \theta \\
&= (x+1)^2 + (y+1)^2 - 2(x+1)(y+1) \cos \theta \\
&= x^2 + y^2 + 2xy \cos \theta + 2(x+y+1)(1 - \cos \theta) \\
&\geq x^2 + y^2 + 2xy \cos \theta &(\cos \theta \leq 1) \\
&= |B'C'|^2.
\end{aligned}
$$

If $|AB|, |AC| \leq 1$, then $B' = C' = A$ and then the claim is trivially true. Suppose $|AB| > 1, |AC| \leq 1$. Now
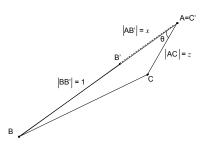

Figure 3: Geometric explanation of Lemma 5.1 when $|AB| > 1, |AC| \leq 1$.

$C' = A$. Let $|AB'| = x, |AC| = z \leq 1$ and θ be the angle at A as shown in Figure 3. Then by the cosine formula,

$$
\begin{aligned}
|BC|^2 &= |AB|^2 + |AC|^2 - 2|AB||AC| \cos \theta \\
&= (x+1)^2 + z^2 - 2(x+1)z \cos \theta \\
&= x^2 + 2x(1 - z \cos \theta) + (z - \cos \theta)^2 + (1 - \cos^2 \theta) \\
&\geq x^2 = |AB'|^2 = |B'C'|^2. &(0 \leq z \leq 1, |\cos \theta| \leq 1)
\end{aligned}
$$

By symmetry, the claim is also true when $|AC| > 1, |AB| \leq 1$. $\qquad \square$

Lemma 5.2. Let $\mathcal{I} = \{(G_1, G_2) : \|G_1 - G_2\|_{\ell_2} \leq 1\}$. Then the ℓ_2-DESCENT update policy is ℓ_2-contractive over the invariant set $\mathcal{I}$.

Proof. Let ϕ denote the ℓ_2-DESCENT policy. Property (2) of Definition 1.3 follows easily because each new user can only change a weighted histogram by an ℓ_2-distance of at most 1.

We will now prove Property (1) of Definition 1.3. Let $(G_1, G_2) \in \mathcal{I}$, i.e., $\|G_1 - G_2\|_{\ell_2} \leq 1$. Let $G_1' = \phi(G_1)$ and $G_2' = \phi(G_2)$. A new user can increase G_1 and G_2 by at most 1 in ℓ_2 distance. Let Γ be the cutoff

parameter in Algorithm 5. Let S be the set of Δ_0 items with the new user, therefore only the items in S will change in G'_1, G'_2. Therefore we can just assume that G_1, G_2 are supported on S for the sake of the analysis. Algorithm 6 moves G_i towards $P = (\Gamma, \Gamma, \ldots, \Gamma)$ by an ℓ_2-distance of 1 (or to P if the distance to P is already lower than 1). We can restrict ourselves to the plane containing G_1, G_2, P (G'_1, G'_2 will also lie on the same plane). Now by Lemma 5.1, $\|G'_1 - G'_2\|_{\ell_2} \leq \|G_1 - G_2\|_{\ell_2}$. $\qquad\square$

We now state a formal theorem which proves $(\varepsilon, \delta) - DP$ of POLICY GAUSSIAN algorithm. We prove it in the supplementary material.

Theorem 5.1. The POLICY GAUSSIAN algorithm (Algorithm 6) is (ε, δ)-DP if $\sigma, \rho_{\mathsf{Gauss}}$ are chosen s.t.

$$\Phi\left(\frac{1}{2\sigma} - \varepsilon\sigma\right) - e^\varepsilon \Phi\left(-\frac{1}{2\sigma} - \varepsilon\sigma\right) \leq \frac{\delta}{2} \text{ and}$$

$$\rho_{\mathsf{Gauss}} \geq \max_{1 \leq t \leq \Delta_0}\left(\frac{1}{\sqrt{t}} + \sigma\Phi^{-1}\left(\left(1 - \frac{\delta}{2}\right)^{1/t}\right)\right).$$

Proof of Theorem 5.1. Suppose D_1 and D_2 are neighboring databases where D_1 has one extra user compared to D_2. Let P and Q denote the distribution of output of the algorithm when the database is D_1 and D_2 respectively. We want to show that $P \approx_{\varepsilon,\delta} Q$. Let E be the event that $A \subset \mathrm{supp}(H_2)$.

Claim 5.1. $P|_E \approx_{\varepsilon,\delta/2} Q$

Proof. Let H_1 and H_2 be the histograms generated by the algorithm from databases D_1 and D_2 respectively. And $\hat{H}_1$ and $\hat{H}_2$ be the histograms obtained by adding $\mathcal{N}(0, \sigma^2)$ noise to each entry of H_1 and H_2 respectively. By the post-processing lemma (Lemma 2.2), it is enough to show that the distributions of the $\hat{H}_1|_{\mathrm{supp}(H_2)}$ and $\hat{H}_2$ are $(\varepsilon, \delta/2)$-close to each other. Because the histogram building algorithm (Algorithm 2) has ℓ_2-sensitivity of at most 1 by hypothesis, $\left\|H_1|_{\mathrm{supp}(H_2)} - H_2\right\|_{\ell_2} \leq 1$. Therefore by properties of Gaussian mechanism (Proposition 2.2), it is enough to choose σ as in the statement of the theorem. $\qquad\square$

By Lemma 2.1, it is enough to show that $P(E) \geq 1 - \delta/2$. Let $T = \mathrm{supp}(H_1) \setminus \mathrm{supp}(H_2)$. Note that $|T| \leq \Delta_0$ and $H_1[u] \leq \frac{1}{\sqrt{|T|}}$ for $u \in T$.

$$
\begin{aligned}
P(\bar{E}) &= \Pr[\exists u \in T \mid \hat{H}_1[u] > \rho_{\mathsf{Gauss}}] \\
&= 1 - \Pr[\forall u \in T \ \ \hat{H}_1[u] \leq \rho_{\mathsf{Gauss}}] \\
&= 1 - \prod_{u \in T} \Pr[\hat{H}_1[u] \leq \rho_{\mathsf{Gauss}}] \\
&= 1 - \prod_{u \in T} \Pr[H_1[u] + X_u \leq \rho_{\mathsf{Gauss}}] & (X_u \sim \mathcal{N}(0, \sigma^2)) \\
&\leq 1 - \prod_{u \in T} \Pr\left[X_u \leq \rho_{\mathsf{Gauss}} - \frac{1}{\sqrt{|T|}}\right] & \left(H_1[u] \leq \frac{1}{\sqrt{|T|}} \text{ for } u \in T\right) \\
&= 1 - \prod_{u \in T} \Phi\left(\frac{\rho_{\mathsf{Gauss}}}{\sigma} - \frac{1}{\sqrt{|T|}}\right)^{|T|} & (3)
\end{aligned}
$$

Thus for

$$\rho_{\mathsf{Gauss}} \geq \max_{1 \leq t \leq \Delta_0}\left(\frac{1}{\sqrt{t}} + \sigma\Phi^{-1}\left(\left(1 - \frac{\delta}{2}\right)^{1/t}\right)\right),$$

we have $P(\bar{E}) \leq \delta/2$. Therefore the DP Set Union algorithm (Algorithm 1) is (ε, δ)-DP. $\qquad\square$

6 Experiments

While the algorithms we described generalize to many domains that involve the release of set union, our experiments will use a natural language dataset. In the context of n-gram release, D is a database of users where each user is associated with 1 or more Reddit posts and W_i is the set of unique n-grams used by each user. The goal is to output as large a subset of n-grams $\cup_i W_i$ as possible while providing (ε, δ)-differential privacy to each user. In our experiments we consider $n = 1$ (unigrams)[3].
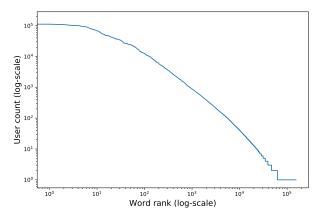
6.1 Dataset

Figure 4: Frequency (i.e. number of users who use the unigram) vs. rank of the unigram (based on frequency) on a log-log scale. This linear relationship shows that the frequency of unigrams among users also follows Zipf's law (power law), i.e., count $\propto 1/\text{rank}^\alpha$ for some constant $\alpha > 0$. The α in this case is ≈ 1.

Our dataset is collected from the subreddit `r/AskReddit`. We take a sample of 15,000 posts from each month between January 2017 and December 2018. We filter out duplicate entries, removed posts, and deleted authors. For text preprocessing, we remove URLs and symbols, lowercase all words, and tokenize using `nltk.word_tokenize`. After preprocessing, we again filter out empty posts to arrive at a dataset of 373,983 posts from 223,388 users.

Similar to other natural language datasets, this corpus follows Zipf's law across users. The frequency of unigrams across users is inversely proportional to the rank of the unigram. Using a log-log scale, the frequency of users for each unigram vs. the rank of the unigram is linear (Figure 4). In other words, the lowest ranked (most common) unigrams are used by almost all users while the highest ranked (least common) unigrams are used by very few users.

Table 1: Percentage of users with unique unigram count of less than or equal to T. The vast majority of users have less than 100 unique unigrams.

THRESHOLD (T)	USERS WITH $|W_i| \leq$ T
1	2.78%
10	29.82%
50	79.16%
100	93.13%
300	99.59%

The distribution of how many unigrams each user uses also follows a long tail distribution. While the top 10 users contribute between 850 and 2000 unique unigrams, most users (93.1%) contribute less than 100

---

[3]The code and dataset used for our experiments are available at `https://github.com/heyyjudes/differentially-private-set-union`

unique unigrams. Table 1 summarizes the percentage of users with a unique vocabulary smaller than each threshold T provided.

6.2 Results

Table 2: Count of unigrams released by various set union algorithms. Results are averaged across 5 shuffles of user order. The best results for each algorithm are in bold. The privacy parameters are $\varepsilon = 3$ and $\delta = \exp(-10)$. The cutoff Γ is calculated using $\alpha = 5$.

Δ_0	1	10	50	100	200	300
Count Laplace	**4484** $\pm$ 32	3666 $\pm$ 7	2199 $\pm$ 8	1502 $\pm$ 14	882 $\pm$ 4	647 $\pm$ 4
Count Gaussian	3179 $\pm$ 15	6616 $\pm$ 18	**6998** $\pm$ 23	6470 $\pm$ 12	5492 $\pm$ 14	4813 $\pm$ 14
Weighted Laplace	**4479** $\pm$ 26	4309 $\pm$ 15	4012 $\pm$ 10	3875 $\pm$ 9	3726 $\pm$ 17	3648 $\pm$ 12
Weighted Gaussian	3194 $\pm$ 11	6591 $\pm$ 18	8570 $\pm$ 14	8904 $\pm$ 24	**8996** $\pm$ 30	8936 $\pm$ 12
Policy Laplace	4482 $\pm$ 21	12840 $\pm$ 28	**15268** $\pm$ 10	14739 $\pm$ 23	14173 $\pm$ 25	13870 $\pm$ 23
Policy Gaussian	3169 $\pm$ 13	11010 $\pm$ 15	16181 $\pm$ 33	16954 $\pm$ 58	**17113** $\pm$ 16	17022 $\pm$ 57

For the problem of outputting the large possible set of unigrams, Table 2 and Figure 5, summarize the performance of DP set union algorithms for different values of Δ_0. The privacy parameters are $\varepsilon = 3$ and $\delta = \exp(-10)$. We compare our algorithms with baseline algorithms: Count Laplace, Count Gaussian, Weighted Laplace, and Weighted Gaussian discussed previously.
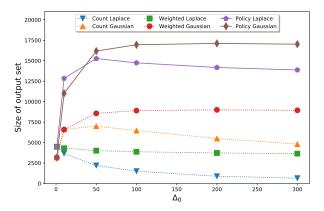

Figure 5: Count of unigrams released by set union algorithms averaged across 5 shuffles of user order. Privacy parameters are $\varepsilon = 3$ and $\delta = \exp(-10)$. The cutoff Γ is calculated using $\alpha = 5$.

Our conclusions are as follows:

- Our new algorithms Policy Laplace and Policy Gaussian output a DP set union that is 2-4 times larger than output of weighted/count based algorithms. This holds for all values of $\varepsilon \geq 1$.

- To put the size of released set in context, we compare our new algorithms against the number of unigrams belonging to at least k users (See Table 3). For Policy Laplace with $\Delta_0 = 100$, the size of the output set covers almost all unigrams (94.8%) when $k = 20$ and surpasses the size of the output set when $k \geq 25$. Policy Gaussian with $\Delta_0 = 100$ covers almost all unigrams (91.8%) when $k = 15$ and surpasses the size of the output set when $k \geq 18$. In other words, our algorithms (with $\varepsilon = 3$ and $\delta = \exp(-10)$) perform better than k-anonymity based algorithms for values of k around 20.

Table 3: This table shows the total number of unigrams that at least k users possess ($|S_k|$) and the percentage coverage of this total by Policy Laplace ($|S_{PL}| = 14739$) and Policy Gaussian ($S_{PG}| = 16954$) for $\Delta_0 = 100$.

k	$|S_k|$	% coverage Policy Laplace	% coverage Policy Gaussian
5	34699	24.5%	48.9%
10	23471	62.8%	72.2%
15	18461	79.8%	91.8%
18	16612	88.7%	102.1%
20	15550	94.8%	109.0%
25	13638	108.1%	124.3%

Table 4: Count of unigrams released Policy Laplace and Policy Gaussian algorithms for single and double passes over users. Results are averaged and rounded across 5 shuffles of user order. The privacy parameters are $\varepsilon = 3$ and $\delta = \exp(-10)$. $\alpha = 2$ is chosen for the threshold parameter. Significant p-values for a two-sided independent t-test are in bold.

	Policy Laplace			Policy Gaussian		
Δ_0	1 Pass	2 Passes	P-val	1 Pass	2 Passes	P-val
1	4236 ± 14	4257 ± 17	0.083	3135 ± 25	3131 ± 20	0.829
10	$\mathbf{12452 \pm 31}$	12389 ± 17	**0.008**	10784 ± 22	10817 ± 54	0.293
50	15056 ± 35	15080 ± 21	0.262	15763 ± 33	15809 ± 45	0.139
100	14562 ± 50	14567 ± 24	0.846	14562 ± 50	14568 ± 24	0.846
200	14005 ± 33	13979 ± 31	0.271	14005 ± 33	13979 ± 31	0.271
300	13702 ± 37	13678 ± 47	0.448	13702 ± 37	13678 ± 47	0.447

6.2.1 Multiple passes through each user

In the experiments described thus far, each user contributes items once within the budget constraints. We also investigate whether the output of set union increases in size when each user contributes the same budget over multiple passes (e.g. user 1 contributes half of their budget each time over 2 passes), we compare Policy Laplace and Policy Gaussian outputs. Table 4 summarizes the results showing that there is not strong evidence suggesting that running multiple passes through the users improves the size of the output set.

6.2.2 Selecting α: parameter to set threshold Γ

Figure 6 shows the number of unigrams released by Policy Laplace and Policy Gaussian for various values of α. We observe that the number of unigrams released increases sharply until $\alpha = 4$, then remains nearly constant and then slowly decreases. This choice of α only affects the policy algorithms since the weighted and count algorithms do not use a threshold.

6.2.3 The effect of ε

We use $\varepsilon = 3$ for the experiments in Table 2. At this value of ε our policy algorithms perform much better than previous count and weighted algorithms. To check whether this result holds with smaller ε, we also run these algorithms on various values of ε. Figure 1 shows that for $\varepsilon \geq 1$ our policy algorithms always perform better.

6.2.4 Selecting hyperparameters while maintaining privacy

As can be seen from Table 2 the Δ_0 resulting in the largest output set varies by algorithm. Since most users in our dataset possess less than 300 unique unigrams, it is not surprising that the largest output set can be
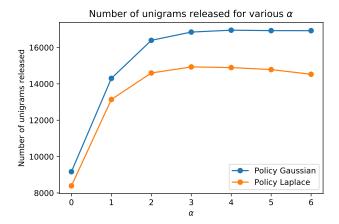
Figure 6: Number of unigrams released for various values of α. The number of unigrams released increases sharply until about $\alpha = 2$, then remains nearly constant and then decreases. Here we fixed $\Delta_0 = 100$ and $\varepsilon = 3$.

achieved with $\Delta_0 < 300$. However, running our algorithms for different values of Δ_0 and selecting the best output will result in a higher value of ε. There are several ways to find the best value of Δ_0 (or any other tunable parameter): 1) using prior knowledge of the data 2) running the algorithms on a small sample of the data to find the best parameters, and discarding that sample. 3) finally, one could also run all the algorithms in parallel and choose the best performing one. Here we will have to account for the loss in privacy budget; see [LT19] for example.

7   Conclusions and open problems

We initiated the study of differentially private set union problem, which has many real-world applications. We designed better algorithms for the problem using the notion of contractive update policy as a guiding principle. One immediate question is to give theoretical guarantees on the size of set union produced by our algorithms. A more interesting and significantly challenging question is to design instance optimal algorithms for the problem. Given the ubiquitous nature of this problem, we believe that it is a worthwhile direction to explore.

References

[Abo16]   John M. Abowd. The challenge of scientific reproducibility and privacy protection for statistical agencies. Technical report, Census Scientific Advisory Committee, 2016.

[AKZ+17]   Brendan Avent, Aleksandra Korolova, David Zeber, Torgeir Hovden, and Benjamin Livshits. Blender: enabling local search with a hybrid differential privacy model. In *Proc. of the 26th USENIX Security Symposium*, pages 747–764, 2017.

[App17]   Differential Privacy Team Apple. Learning with privacy at scale. Technical report, Apple, 2017.

[BEM+17]   Andrea Bittau, Úlfar Erlingsson, Petros Maniatis, Ilya Mironov, Ananth Raghunathan, David Lie, Mitch Rudominer, Ushasree Kode, Julien Tinnes, and Bernhard Seefeld. Prochlo: Strong privacy for analytics in the crowd. In *Proceedings of the 26th Symposium on Operating Systems Principles*, SOSP '17, pages 441–459, 2017.

[BW18] Borja Balle and Yu-Xiang Wang. Improving the gaussian mechanism for differential privacy: Analytical calibration and optimal denoising. In *International Conference on Machine Learning*, pages 403–412, 2018.

[CLB+19] Mia Xu Chen, Benjamin N. Lee, Gagan Bansal, Yuan Cao, Shuyuan Zhang, Justin Lu, Jackie Tsay, Yinan Wang, Andrew M. Dai, Zhifeng Chen, and et al. Gmail smart compose: Real-time assisted writing. In *Proceedings of the 25th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, page 2287–2295, 2019.

[CLE+19] Nicholas Carlini, Chang Liu, Úlfar Erlingsson, Jernej Kos, and Dawn Song. The secret sharer: Evaluating and testing unintended memorization in neural networks. In *28th {USENIX} Security Symposium ({USENIX} Security 19)*, pages 267–284, 2019.

[DBS19] Budhaditya Deb, Peter Bailey, and Milad Shokouhi. Diversifying reply suggestions using a matching-conditional variational autoencoder. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 2 (Industry Papers)*, Minneapolis, Minnesota, June 2019. Association for Computational Linguistics.

[DKY17] Bolin Ding, Janardhan Kulkarni, and Sergey Yekhanin. Collecting telemetry data privately. In *Advances in Neural Information Processing Systems*, pages 3574–3583, 2017.

[DMNS06] Cynthia Dwork, Frank McSherry, Kobbi Nissim, and Adam Smith. Calibrating noise to sensitivity in private data analysis. In *Theory of Cryptography Conference*, pages 265–284. Springer, 2006.

[DR14a] Cynthia Dwork and Aaron Roth. The algorithmic foundations of differential privacy. *Foundations and Trends in Theoretical Computer Science*, 9(3–4):211–407, 2014.

[DR+14b] Cynthia Dwork, Aaron Roth, et al. The algorithmic foundations of differential privacy. *Foundations and Trends® in Theoretical Computer Science*, 9(3–4):211–407, 2014.

[EPK14] Úlfar Erlingsson, Vasyl Pihur, and Aleksandra Korolova. Rappor: Randomized aggregatable privacy-preserving ordinal response. In *Proceedings of the 2014 ACM SIGSAC conference on computer and communications security*, pages 1054–1067. ACM, 2014.

[FMTT18] Vitaly Feldman, Ilya Mironov, Kunal Talwar, and Abhradeep Thakurta. Privacy amplification by iteration. In Mikkel Thorup, editor, *59th IEEE Annual Symposium on Foundations of Computer Science, FOCS 2018, Paris, France, October 7-9, 2018*, pages 521–532. IEEE Computer Society, 2018.

[HLLC14] Baotian Hu, Zhengdong Lu, Hang Li, and Qingcai Chen. Convolutional neural network architectures for matching natural language sentences. In *Advances in neural information processing systems*, pages 2042–2050, 2014.

[KCK+18] Yu-Hsuan Kuo, Cho-Chun Chiu, Daniel Kifer, Michael Hay, and Ashwin Machanavajjhala. Differentially private hierarchical group size estimation. *arXiv preprint arXiv:1804.00370*, 2018.

[KKMN09] Aleksandra Korolova, Krishnaram Kenthapadi, Nina Mishra, and Alexandros Ntoulas. Releasing search queries and clicks privately. In *Proceedings of the 18th international conference on World wide web*, pages 171–180, 2009.

[KKR+16] Anjuli Kannan, Karol Kurach, Sujith Ravi, Tobias Kaufmann, Andrew Tomkins, Balint Miklos, Greg Corrado, Laszlo Lukacs, Marina Ganea, Peter Young, et al. Smart reply: Automated response suggestion for email. In *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, pages 955–964, 2016.

[LT19] Jingcheng Liu and Kunal Talwar. Private selection from private candidates. In Moses Charikar and Edith Cohen, editors, *Proceedings of the 51st Annual ACM SIGACT Symposium on Theory of Computing, STOC 2019, Phoenix, AZ, USA, June 23-26, 2019*, pages 298–309. ACM, 2019.

[MT07] Frank McSherry and Kunal Talwar. Mechanism design via differential privacy. In *48th Annual IEEE Symposium on Foundations of Computer Science (FOCS 2007), October 20-23, 2007, Providence, RI, USA, Proceedings*, pages 94–103. IEEE Computer Society, 2007.

[Vad17] Salil Vadhan. The complexity of differential privacy. In *Tutorials on the Foundations of Cryptography*, pages 347–450. Springer, 2017.

[WZL+20] Royce Wilson, Celia Yuxin Zhang, William Lam, Damien Desfontaines, Daniel Simmons-Marengo, and Bryant Gipson. Differentially private SQL with bounded user contribution. 2020.

A   Bounded Sensitivity implies DP (Proof of Theorem 1.2)

In this section, we will prove a formal version of Theorem 1.2, i.e., if the histogram output by Algorithm 2 has bounded ℓ_p-sensitivity (for $p \in \{1, 2\}$), then by adding appropriate noise and setting an appropriate threshold, Algorithm 1 for DP set union can be made differentially private. The lower bounds on the threshold (ρ) that we obtain in this generality are only slightly worse compared to the corresponding bounds in Theorems 4.1 and 5.1.

Theorem A.1. Suppose the histogram output by Algorithm 2 has ℓ_1-sensitivity 1. Then Algorithm 1 is (ε, δ)-DP when the Noise distribution is $\mathsf{Lap}(0, \lambda)$ where $\lambda = 1/\varepsilon$ and the threshold

$$\rho \geq \max_{1 \leq t \leq \Delta_0} 1 + \frac{1}{\varepsilon} \log \left(\frac{1}{2 \left(1 - (1 - \delta)^{1/t}\right)} \right).$$

Proof. Proof of Theorem A.1 is extremely similar to the proof of Theorem 4.1. The only place where it differs is in Equation (2) where we bound $H_1[u] \leq 1$ instead of $H_1[u] \leq 1/|T|$. $\qquad \square$

Theorem A.2. Suppose the histogram output by Algorithm 2 has ℓ_2-sensitivity 1. Then Algorithm 1 is (ε, δ)-DP when the Noise distribution is $\mathcal{N}(0, \sigma^2)$ where σ and the threshold ρ are chosen s.t.

$$\Phi \left(\frac{1}{2\sigma} - \varepsilon\sigma \right) - e^\varepsilon \Phi \left(-\frac{1}{2\sigma} - \varepsilon\sigma \right) \leq \frac{\delta}{2} \text{ and}$$

$$\rho \geq \max_{1 \leq t \leq \Delta_0} \left(1 + \sigma \Phi^{-1} \left(\left(1 - \frac{\delta}{2}\right)^{1/t} \right) \right).$$

Proof. Proof of Theorem A.2 is extremely similar to the proof of Theorem 5.1. The only place where it differes is in Equation (3) where we bound $H_1[u] \leq 1$ instead of $H_1[u] \leq 1/\sqrt{|T|}$. $\qquad \square$

B   Weighted Laplace and Gaussian algorithms

B.1   Weighted Laplace

Algorithm 7 Laplace weighted update

Input: H: Current histogram
W: A subset of U of size at most Δ_0
Output: H: Updated histogram
for u in W **do**
$\quad$ H$[u] \leftarrow$ H$[u] + \frac{1}{|W|}$
end for

Theorem B.1. The Weighted Laplace algorithm (Algorithm 7) is (ε, δ)-DP when

$$\rho_{\mathsf{Lap}} \geq \max_{1 \leq t \leq \Delta_0} \frac{1}{t} + \frac{1}{\varepsilon} \log \left(\frac{1}{2 \left(1 - (1 - \delta)^{1/t}\right)} \right).$$

Proof. Proof is exactly the same as that of Theorem 4.1. $\qquad \square$

B.2   Weighted Gaussian

Algorithm 8 GAUSSIAN weighted update

Input: H: Current histogram
W: A subset of U of size at most Δ_0
Output: H: Updated histogram
for u in W **do**
    $H[u] \leftarrow H[u] + \sqrt{\frac{1}{|W|}}$
end for

Theorem B.2. The WEIGHTED GAUSSIAN algorithm (Algorithm 8) is (ε, δ)-DP if $\sigma, \rho_{\mathsf{Gauss}}$ are chosen s.t.

$$\Phi\left(\frac{1}{2\sigma} - \varepsilon\sigma\right) - e^\varepsilon \Phi\left(-\frac{1}{2\sigma} - \varepsilon\sigma\right) \leq \frac{\delta}{2} \text{ and}$$

$$\rho_{\mathsf{Gauss}} \geq \max_{1 \leq t \leq \Delta_0} \left(\frac{1}{\sqrt{t}} + \sigma\Phi^{-1}\left(\left(1 - \frac{\delta}{2}\right)^{1/t}\right)\right).$$

Proof. Proof is exactly the same as that of Theorem 5.1. □
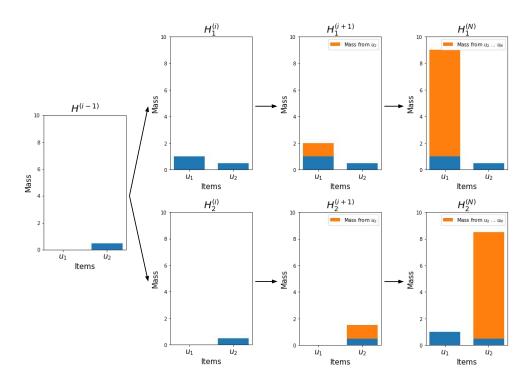
C   Greedy policy

Figure 7: Visualization of greedy update example where the final l_1 sensitivity is larger than 1.

In this section, we give a simple counter example to illustrate how the sensitivity of a greedy policy algorithm can be unbounded.

Suppose there are N users, let u_1 and u_2 be two items in the universe. We will denote the weight of item u after user i's contribution as $H^{(i)}[u]$. Suppose user i has only item u_1 while users $i+1, i+2, \ldots, N$ have both items. Let H_1 be the histogram generated with all N users while H_2 be the histogram generated without user i. Let $\Delta_0 = 2$ and $H^{(i-1)}[u_1] < H^{(i-1)}[u_2] < 1 + H^{(i-1)}[u_1]$. According to the greedy update

Algorithm 9 GREEDY POLICY update
___
Input: H: Current histogram
W: A subset of U of size at most Δ_0
Γ: cutoff parameter
Output: H: Updated histogram
// Build cost dictionary G
$G = \{\}$ // Empty dictionary
for $u \in W$ **do**
  if $H[u] < \Gamma$ **then**
    // Gap to cutoff for items below cutoff Γ
    $G[u] \leftarrow \Gamma - H[u]$
  end if
end for
budget $\leftarrow 1$ // Each user gets a total budget of 1
// Sort in increasing order of the gap $\Gamma - H[u]$
$G \leftarrow \text{sort}(G)$
// Let $u_1, u_2, \ldots, u_{|G|}$ be the sorted order
for $j = 1$ to $|G|$ **do**
  if $G[u_j] \leq$ budget **then**
    $H[u_j] \leftarrow H[u_j] + G[u_j]$
    budget $\leftarrow$ budget - $G[u_j]$
  else
    $H[u_j] \leftarrow H[u_j] +$ budget
    break
  end if
end for
___

described in Algorithm 9, in H_1, user i will add weight 1 to u_1 and users $i+1, i+2, \ldots, N$ will also to u_1 since $H^{(i)}[u_1] > H^{(i)}[u_2]$. In H_2, users $i+1, i+2, \ldots, N$ will add to u_2 since $H^{(i-1)}[u_1] < H^{(i-1)}[u_2]$. This process is described in figure 7. Therefore the ℓ_1-sensitivity of the histogram built using Greedy Policy update (Algorithm 9) can be $\Omega(\Gamma, N)$.